

Secure Programming I

Davide Balzarotti
davide@iseclab.org

Overview


- Parallel execution of tasks
 - multi-process or multi-threaded environment
 - tasks can interact with each other
- Interaction
 - shared memory (or address space)
 - file system
 - signals
- Results of tasks depends on relative timing of events
 - Indeterministic behavior

Race Conditions

- Race conditions
 - alternative term for indeterministic behavior
 - often a robustness issue
 - but also many important security implications
- Assumption needs to hold for some time for correct behavior,
 - but assumption can be violated
- Time window when assumption can be violated
 - window of vulnerability


Race Conditions

- Programmer views a set of operations as atomic
- In reality, atomicity is not enforced
- Attacker can take advantage of this discrepancy



Race Conditions

- Programmer views a set of operations as atomic
- In reality, atomicity is not enforced
- Attacker can take advantage of this discrepancy



Shared Memory

- Sharing of memory between tasks can lead to races
 - Threads share the entire memory space
 - Processes may share memory mapped regions
- Use synchronization primitives:
 - locking, semaphores
 - Java:
 - `synchronized` classes and methods (Monitor model)
 - Atomic types (`java.util.concurrent.atomic.AtomicInteger`, etc)
- Avoid shared memory:
 - use message-passing model
 - still need to get the synchronization right!

Shared Memory

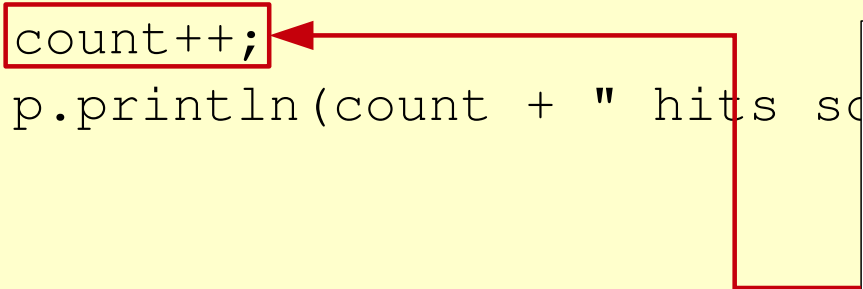
(trivial) example:

```
public class Counter extends HttpServlet {
    int count = 0;
    public void doGet(HttpServletRequest in,
                      HttpServletResponse out)
    {
        out.setContentType("text/plain");
        PrintWriter p = out.getWriter();
        count++;
        p.println(count + " hits so far!");
    }
}
```

Shared Memory

(trivial) example:

```
public class Counter extends HttpServlet {  
    int count = 0;  
    public void doGet(HttpServletRequest in,  
                      HttpServletResponse out)  
    {  
        out.setContentType("text/plain");  
        PrintWriter p = out.getWriter();  
        count++;  
        p.println(count + " hits so far");  
    }  
}
```



- Looks atomic (1 line of code!)
 - It's not!
- Simple race:
 - 2 threads read count
 - both write count+1
 - missed 1 increment


Race Conditions

Sequence of operations (A,B):

- Is not atomic
- can be interrupted at any time for arbitrary amounts of time



Programmer



Scheduler can interrupt a process at
any time

- Can happen between A and B


Race Conditions

Sequence of operations (A,B):

- Is not atomic
- can be interrupted at any time for arbitrary amounts of time



Programmer



Scheduler can interrupt a process at **any time**

- Can happen between A and B
- Much more likely if there is a **blocking system call** in between


Window of Vulnerability

- Things go wrong if **C** happens between t_A and t_B .



Programmer

- (t_A, t_B) is the window of vulnerability



Attacker

Race Conditions

- Window of vulnerability can be very short
 - race condition problems are difficult to find with testing
 - difficult to reproduce and debug
- Myths about race conditions
 - *"races are hard to exploit"*
 - *"races cannot be exploited reliably"*
 - *"only 1 chance in 10000 that the attack will work!"*
- Attackers can often find ways to beat the odds!

Beating the Odds

- Can the attacker try the exploit 1 million times?
 - if yes, and the odds are 1 to 10000, then he has a reliable exploit
- Attacker can try to slow down the victim machine/process to improve the odds
 - high load
 - computational complexity attacks
- Attacker can run the attack many times in parallel to increase the probability that the attacking process will be scheduled by the processor at the right moment

Time-of-Check, Time-of-Use (TOCTOU)

- Time-of-Check, Time-of-Use (TOCTOU)
 - common race condition problem
- Problem:
 - **Time-Of-Check** (t_A): validity of assumption X on entity E is checked
 - **Time-Of-Use** (t_B): assuming X is still valid, E is used
 - **Time-Of-Attack** (t_C): assumption X is invalidated
 - $t_A < t_C < t_B$
- Program has to execute with elevated privilege
 - otherwise, attacker races for his own privileges

Time-of-Check, Time-of-Use

- Steps to access a resource
 - obtain reference to resource
 - query resource to obtain characteristics
 - analyze query results
 - if resource is fit, access it
- Often occurs in Unix file system accesses
 - check permissions for a certain file name (e.g., using `access(2)`)
 - open the file, using the file name (e.g., using `fopen(3)`)
 - four levels of indirection (symbolic link - hard link - inode - file descriptor)

access/open Race


- Case study: setuid program
- `man 2 access`:
"The check is done using the calling process's real UID and GID, rather than the effective IDs as is done when actually attempting an operation (e.g., `open(2)`) on the file. This allows set-user-ID programs to easily determine the invoking user's authority."

```
/* access returns 0 on success */  
if(!access(file, W_OK)) {  
    f = fopen(file, "wb+");  
    write_to_file(f);  
} else {  
    fprintf(stderr, "Permission denied,  
                  cannot open %s.\n", file);  
}
```


access/open Race

```
/* access returns 0 on success */  
if(!access(file, W_OK)) {  
    f = fopen(file, "wb+");  
    write_to_file(f);  
} else {  
    ...  
}
```


\$ touch dummy; ln -s dummy pointer



access/open Race

```
/* access returns 0 on success */
if(!access(file, W_OK)) {
    f = fopen(file, "wb+");
    write_to_file(f);
} else {
    ...
}
```

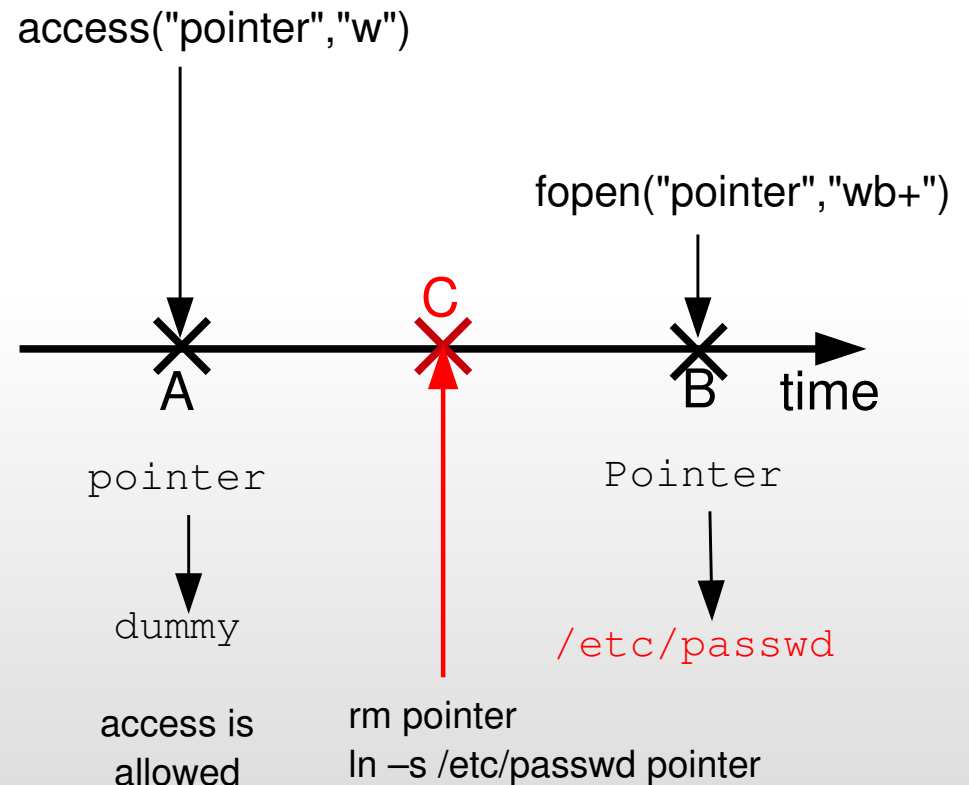
```
$ touch dummy; ln -s dummy pointer
$ rm pointer; ln -s /etc/passwd pointer
```



access/open Race

```
/* access returns 0 on success */  
if(!access(file, W_OK)) {  
    f = fopen(file, "wb+");  
    write_to_file(f);  
} else {  
    ...  
}
```

```
$ touch dummy; ln -s dummy pointer  
$ rm pointer; ln -s /etc/passwd pointer
```

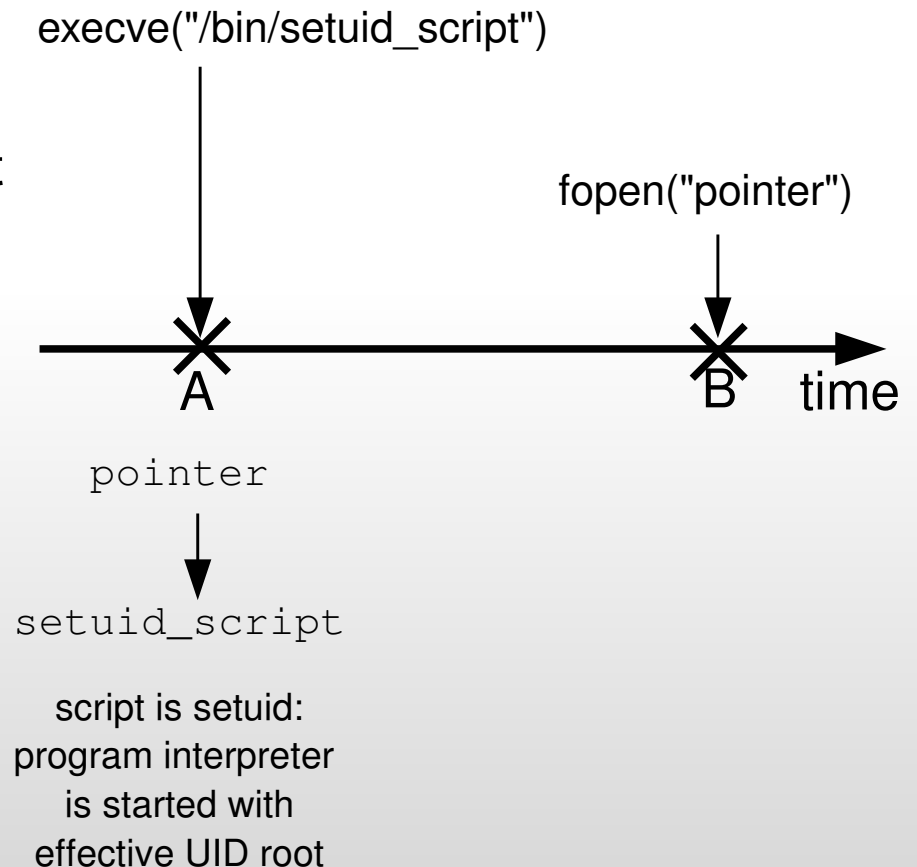


Script execve Race

- Filename redirection
 - soft links again
- Setuid Scripts
 - `execve()` system call invokes `setuid()` call prior to executing program
 - A: program is a script, so command interpreter is loaded first
 - B: program interpreter (with root privileges) is invoked on script name
 - attacker can replace script content between step A and B
- Setuid not allowed on scripts on most platforms!
 - although there are some work-arounds

script execve Race

- A: program interpreter is started (with root privilege)
 - e.g: /bin/sh, /usr/bin/python,
- B: program interpreter opens script pointed to by "pointer"
- Interpreter runs the script



- Attack:

```
$ ln -s /bin/setuid_script pointer
```

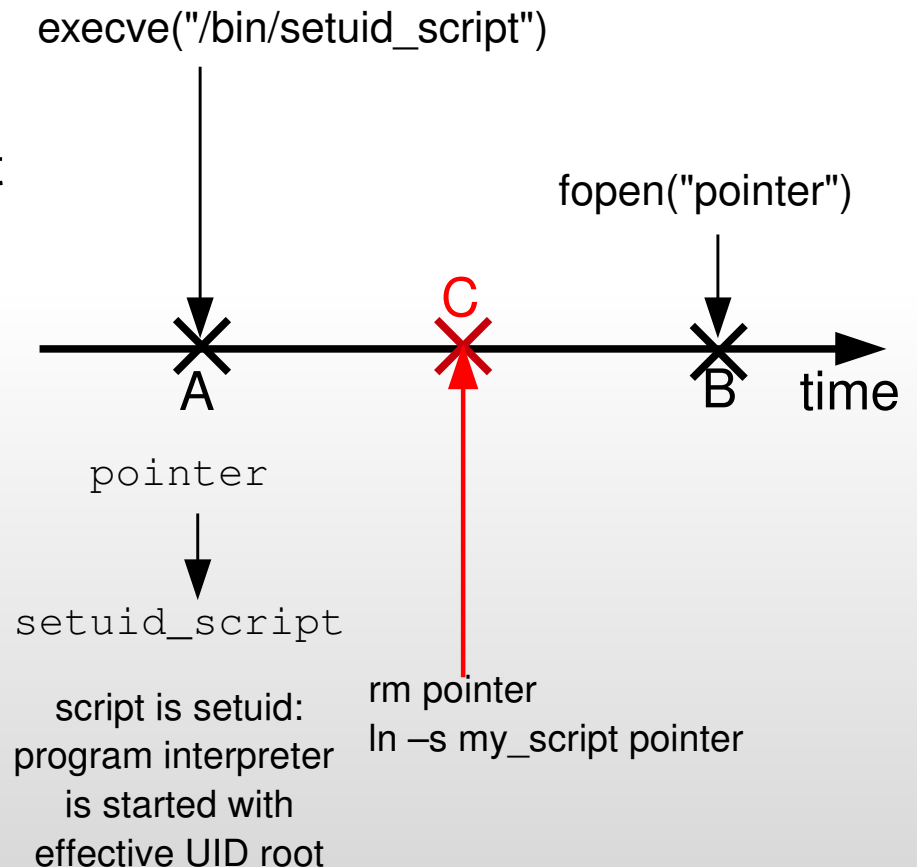
script execve Race

- A: program interpreter is started (with root privilege)
 - e.g: /bin/sh, /usr/bin/python,
- B: program interpreter opens script pointed to by "pointer"
- Interpreter runs the script

- Attack:

```
$ ln -s /bin/setuid_script pointer
```

```
$ rm pointer; ln -s my_script pointer
```



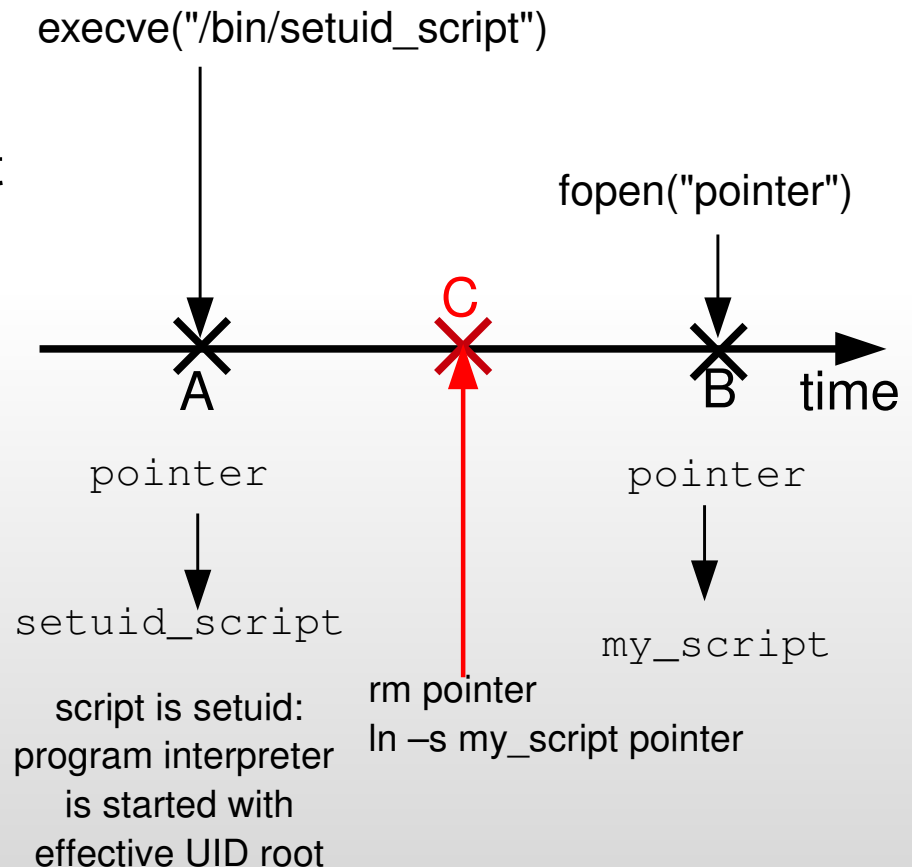
script execve Race

- A: program interpreter is started (with root privilege)
 - e.g: /bin/sh, /usr/bin/python,
- B: program interpreter opens script pointed to by "pointer"
- Interpreter runs the script

- Attack:

```
$ ln -s /bin/setuid_script pointer
```

```
$ rm pointer; ln -s my_script pointer
```



Directory operations

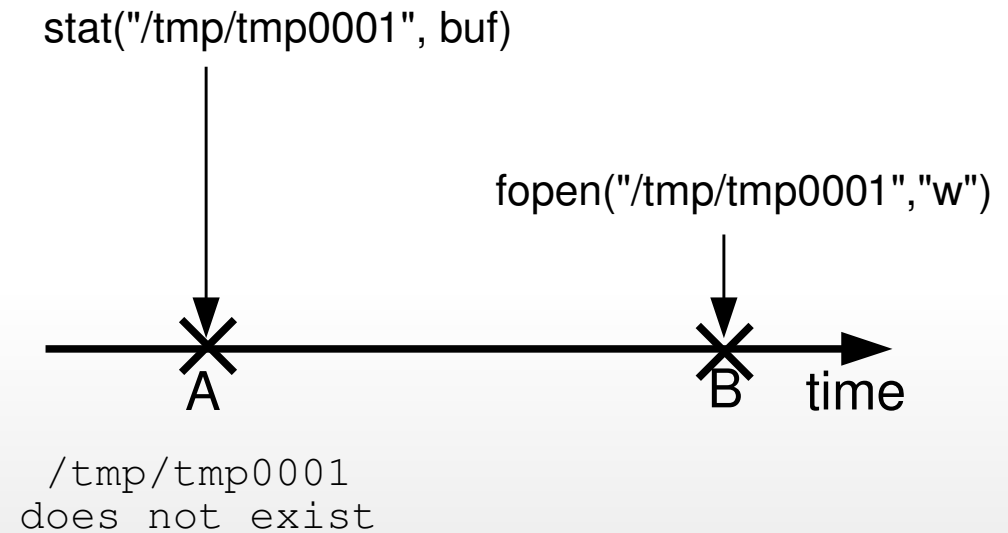
- “`rm -r`” race
 - `rm` can remove directory trees, traverses directories depth-first
 - issues `chdir("..")` to go one level up after removing a directory branch
 - by relocating subdirectory to another directory (while `rm -r` is running!), arbitrary files can be deleted

Races on temporary files

- Similar issues as with regular files
 - commonly opened in /tmp or /var/tmp
 - creating files in /tmp requires no special permissions
 - often guessable file name
- (One) Attack:
 - guess the tmp file name: "/tmp/tmp0001"
 - `ln -s /etc/target /tmp/tmp0001`
 - victim program will create file /etc/target for you, when it tries to create the temporary file!
 - if first guess doesn't work, try 1 million times


Races on temporary files

- A: program checks if file `"/tmp/tmp0001"` already exists
- B: program creates file `"/tmp/tmp0001"`



Races on temporary files

- A: program checks if file `/tmp/tmp0001` already exists
- B: program creates file `/tmp/tmp0001`

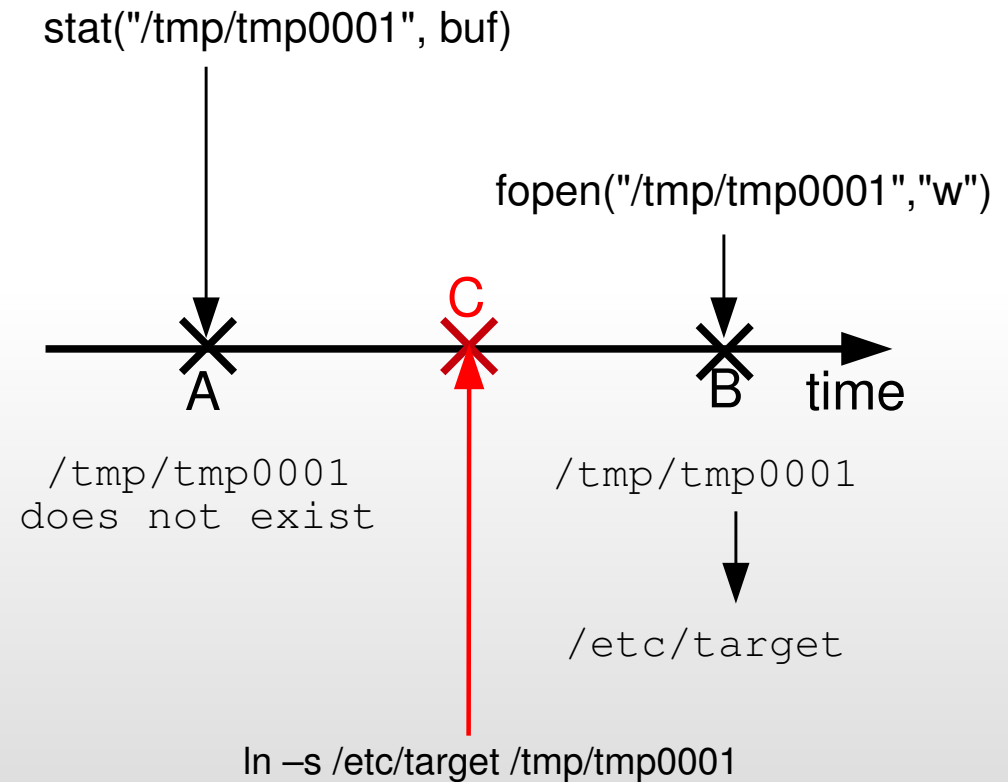


- Attack:

`$ ln -s /etc/target /tmp/tmp0001`

Races on temporary files

- A: program checks if file `/tmp/tmp0001` already exists
- B: program creates file `/tmp/tmp0001`
 - `/etc/target` is created!



- Attack:

\$ `ln -s /etc/target /tmp/tmp0001`

Races on temporary files

- Temp Cleaners
 - programs that clean “old” temporary files from temp directories
 - first `lstat(2)` file, then use `unlink(2)` to remove files
- attack: arbitrary file deletion
 - race condition when attacker replaces file (softlink) between `lstat(2)` and `unlink(2)`
- attack: delete temporary file too early
 - delay program long enough until temp cleaner removes active temporary file

Races on temporary files

- “Secure” procedure
 - pick hard to guess filename (randomize part of name)
 - set umask appropriately (0066 is usually good)
 - **atomically** test for existence AND create the file
 - use `open(2)` `O_CREAT|O_EXCL` to create the file, opening it in the proper mode
 - if file exists, `fopen` will fail: try again with another file name (in a loop)
 - delete the file immediately using `unlink(2)`
 - until the file descriptor is open is still usable by the program
 - perform reads, writes, and seeks on the file as necessary
 - finally, close the file: it is automatically deleted

Races on temporary files

- umask issues
 - if all users have read access, can lead to leak of private data
 - if all users have write access, can lead to data tampering
 - programs treat their temporary files as trusted
 - they may not validate input from them
 - maybe I can find a vulnerability in the program if I can tamper with its temporary files
- Use library functions to create temporary files
 - don't roll your own implementation!
- Some library functions are insecure
 - `mktemp(3)` is not secure, use `mkstemp(3)` instead
 - old versions of `mkstemp(3)` did not set umask correctly

More examples

- File meta-information
 - `chown(2)` and `chmod(2)` are unsafe because they operate on file names
 - use `fchown(2)` and `fchmod(2)` that use file descriptors
- Logging/Crash reporting
 - example: Joe Editor vulnerability
 - when joe crashes (e.g., segmentation fault, xterm crashes) unconditionally append open buffers to local DEADJOE file
 - DEADJOE could be symbolic link to security-relevant file

More examples

- SQL select before insert
 - use select to check if a certain element already exists
 - when select returns no results, insert a (unique) element
- Race condition:
 - 2 processes may do this at the same time, leading to 2 insertions
- Countermeasures
 - locking
 - primary keys: use a single **atomic** insert. It will fail if key already exists

Computational Complexity Attacks

Beating the odds

- Window of vulnerability can be short
- Attacker can try to make the program run more slowly
- Computational complexity attacks
 - many algorithms are fast on average
 - ...but are slow in some corner cases?
- Example: filename lookups
 - deeply nested directory structure
 - chain of symbolic links
 - looking up the file in the FS will take longer!

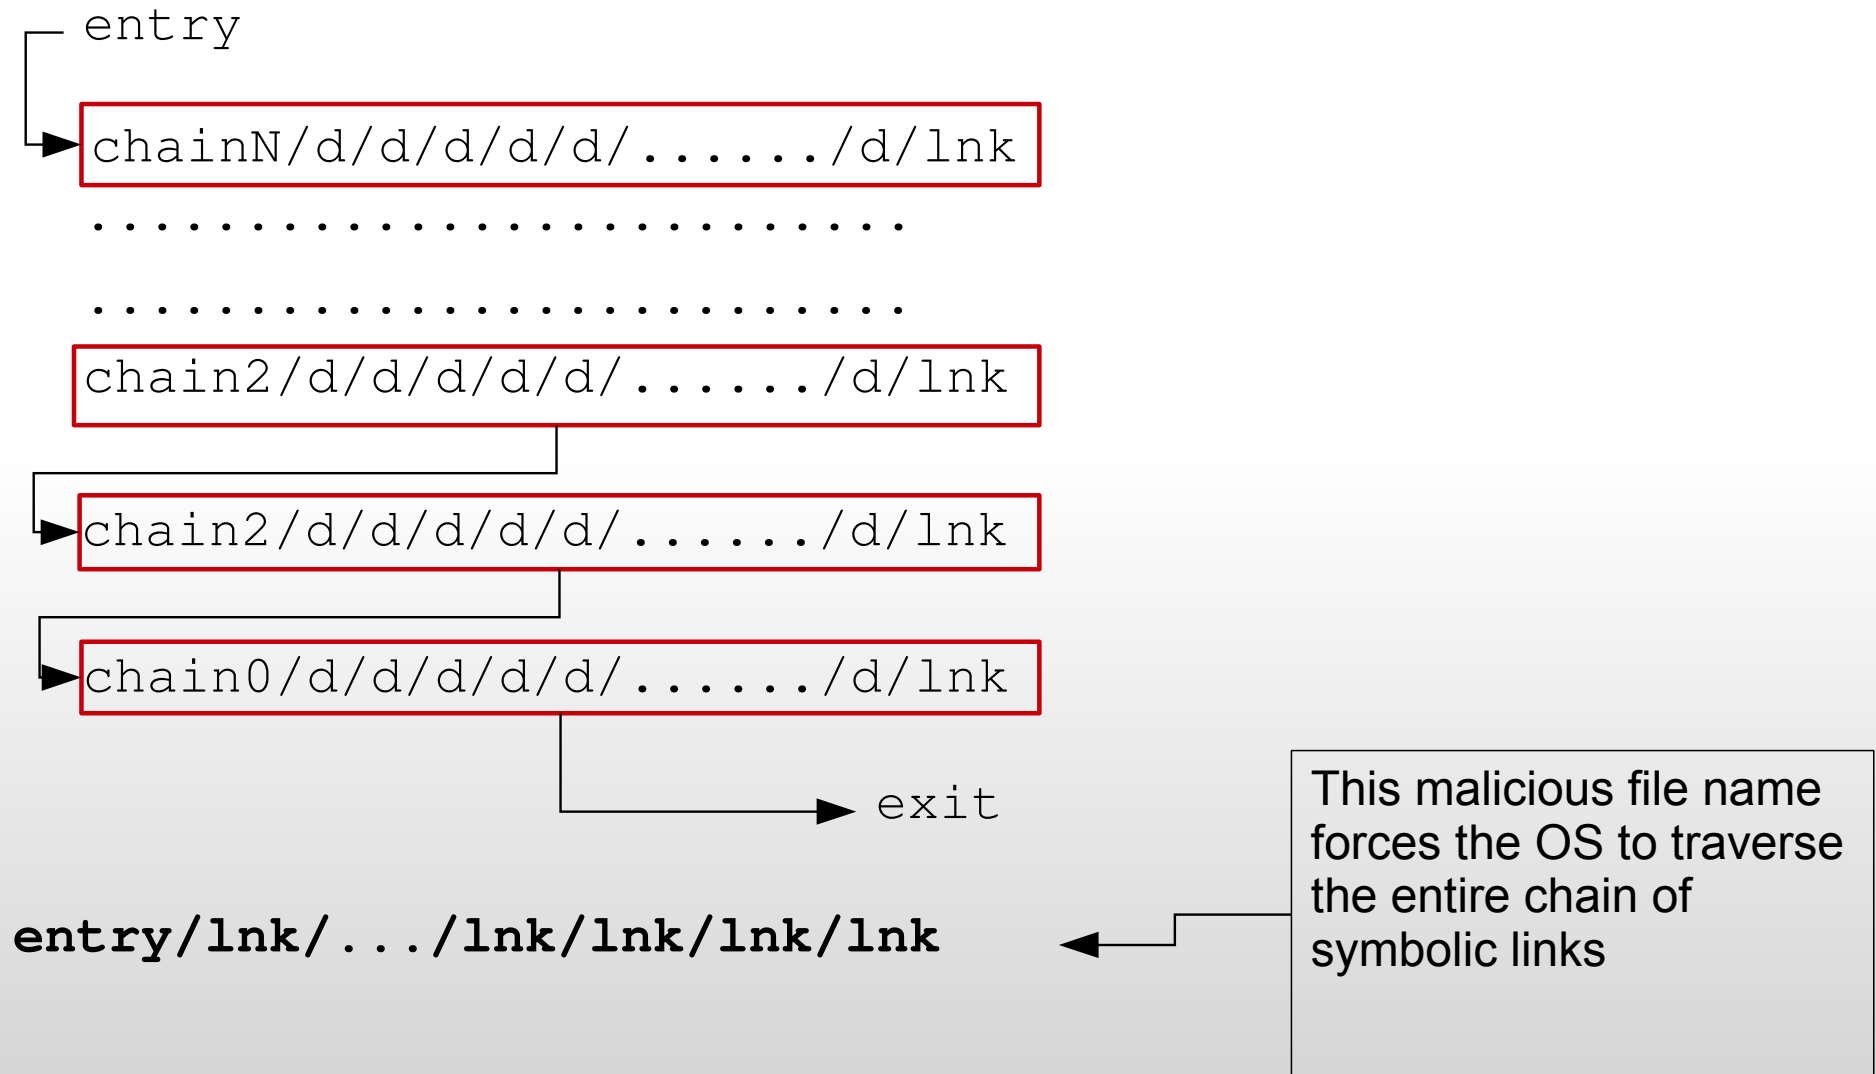
Slow file lookups

- Deeply nested directory structure:
 - `d/d/d/d/d/d/d/...../d/file.txt`
- To resolve this file name, the OS must:
 - look for directory named d in current working directory
 - look for directory named d in that directory
 - ...
 - look for file named file.txt in final directory
- Limit to length of a file name:
 - MAXPATHLENGTH (4096 on my linux)
 - Max depth of ~2000

File System Maze

- Combine deeply nested directory structure with chain of symbolic links
 - MAXPATHLENGTH limits length of file parameter to a single system call (e.g, `open`, `access`)
 - But parts of a file name can themselves be links
 - Length of link chain limited by kernel parameter
 - 40 on my linux box
- Total file system lookups:
 - follow 40 chains...
 - ...each with 2000 nested directories
 - **80000** lookups!

File System Maze



Mini-Demo

```
pmilani@rorschach:~$time cat  
entry/lnk/lnk/lnk/lnk/lnk/lnk/  
/lnk/lnk/lnk/lnk/lnk/lnk/lnk/  
lnk/lnk/lnk/lnk/lnk/lnk/lnk/l  
nk/lnk/lnk/lnk/lnk/lnk/lnk/ln  
k/lnk/lnk/lnk/lnk/lnk/lnk/lnk
```

CONTENT OF FILE

```
real 7m22.498s
```

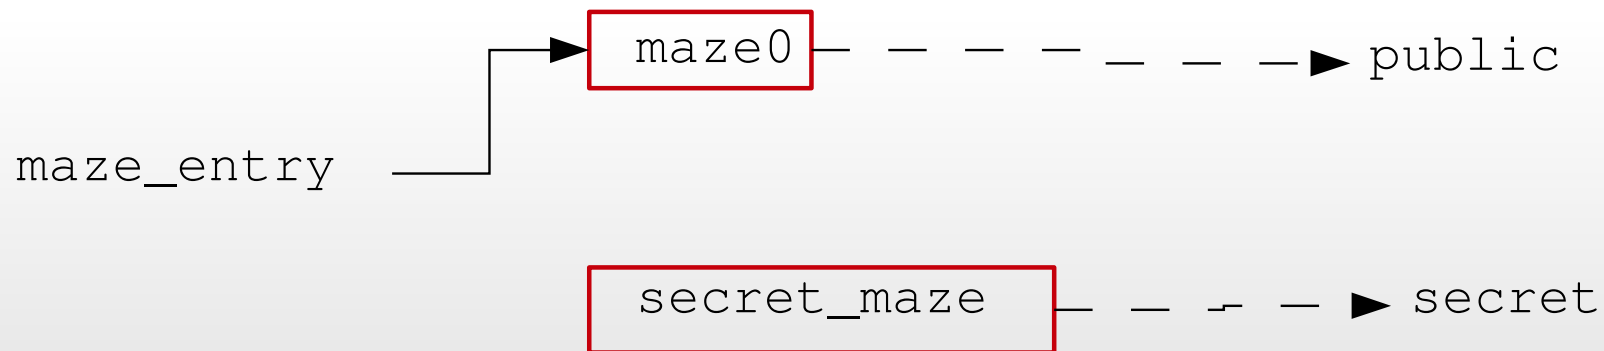
```
user 0m0.000s
```

```
sys 0m4.964s
```

- In my tests, traversing the maze sometimes took over 7 minutes!
 - while keeping the disk busy with other operations
 - with the directories in the maze not already cached by the OS
- Reliably over half a second
- Plenty of time to exploit a race condition

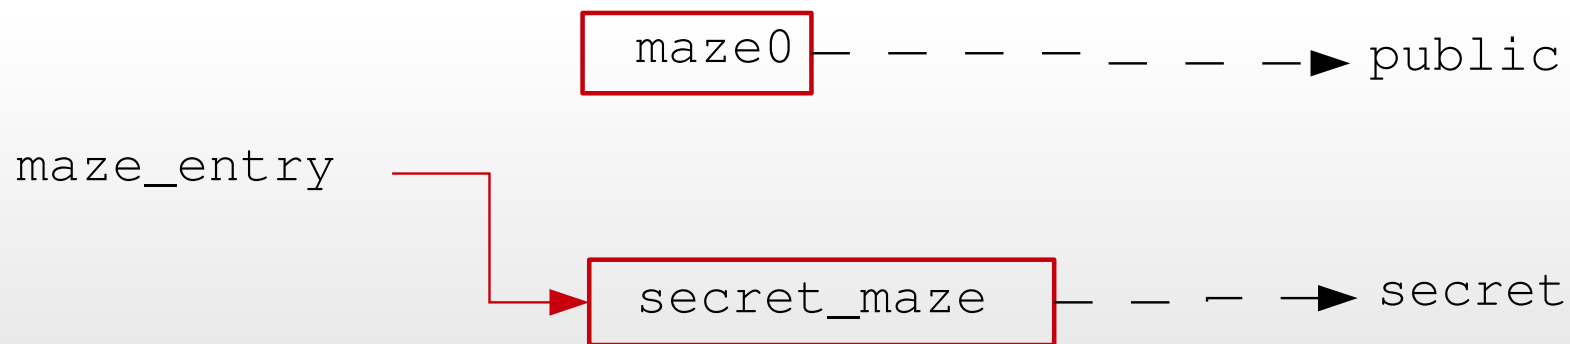
access/open Race Demo

- `suid_cat`:
 - vulnerable program: setuid version of cat utility
 - uses `access` to check if it can open a file
- multiple chains of symbolic links:



access/open Race Demo

- `suid_cat`:
 - vulnerable program: setuid version of cat utility
 - uses `access` to check if it can open a file
- multiple chains of symbolic links:



- Change the `maze_entry` while `suid_cat` is running!

Computational Complexity Attacks

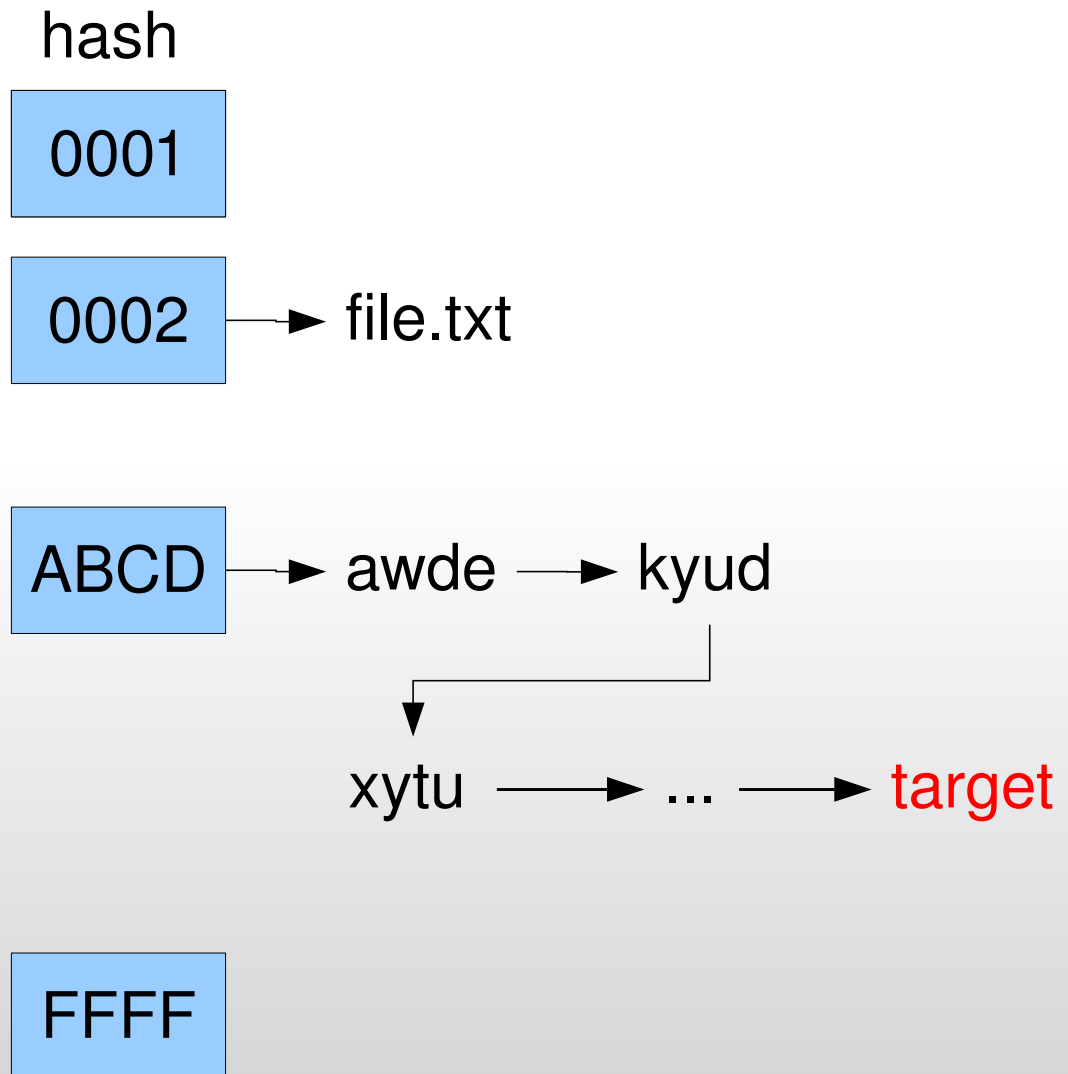
- Exploit worst-case performance of an algorithm
- Example: file lookup (again)
- How does OS store mapping between file names and inodes in a directory?
 - linked list or array?
 - what if there are 100000 files in directory?
 - too slow in practice
 - hash table!
 - good average performance
 - bad worst-case performance
 - can an attacker exploit this?

Hash Table

- Store objects in a number of buckets
- Each bucket is a linked list
- Choose bucket for an object X based on hash function $h(X)$
- Accessing an item in a hash table of N elements is $O(1)$ most of the time
 - because most buckets hold 1 or 0 elements
- Worst case complexity is $O(N)$!
 - if all objects have the same hash
- Worst case does not occur accidentally (very unlikely)
- Attacker can make worst case happen

Computational Complexity Attacks

- File entries in a directory typically stored in a hash table
- Hash tables are slow when there are many entries in the same bucket
- Create 10000 files with same hash!



Detection and Prevention

Prevention


- Do not assume you are safe from race conditions just because:
 - Window of opportunity is short
 - Attacker may well be able to make it bigger!
 - Success is unlikely
 - Attacker may be able to try 1 million times!
 - Now that most of the computers have multiple cores, attacks are even easier

Preventing Race Conditions

- operate on file descriptors
 - not on file names (as much as possible)
- do not check access by yourself
 - (i.e., no use of `access(2)`)
- drop privileges instead and let the file system do the job
- Use the `O_CREAT | O_EXCL` flags to create a new file with `open(2)`
 - and be prepared to have the open call fail

Avoiding the access/open Race


```
ruid=getuid();  
euid=geteuid();  
/* drop privileges */  
seteuid(ruid);  
int fd = open(file, O_CREAT|O_EXCL);  
if(fd!=-1) {  
    write_to_file(fd);  
} else {  
    fprintf(stderr, "Permission denied");  
}  
seteuid(euid);
```



- when acting on behalf of the user, assume his identity
 - let the operating system check permissions!
- this code snippet still has 2 potential vulnerabilities

Avoiding the access/open Race

```
ruid=getuid();  
euid=geteuid();  
/* drop privileges */  
seteuid(ruid);  
int fd = open(file, O_CREAT|O_EXCL);  
if(fd!=-1) {  
    write_to_file(fd);  
} else {  
    fprintf(stderr, "Permission denied");  
}  
seteuid(euid);
```



- check seteuid for errors
 - if seteuid fails, your effective UID is unchanged (you are still root!)
- also drop group privileges with **setegid()**

Prevention

- Some calls require file names

```
link(), mkdir(), mknod(), rmdir(),  
symlink(), unlink()
```

- especially `unlink(2)` is troublesome

- Secure File Access

- create “secure” directory
- directory only write and executable by UID of process
- check that no parent directory can be modified by attacker
- walk up directory tree
checking for permissions and links at each step

Locking

- Ensures exclusive access to a certain resource
 - Used to avoid accidental race conditions
 - advisory locking (processes need to cooperate)
 - not mandatory, therefore not secure
- Often, files are used for locking
 - portable (files can be created nearly everywhere)
 - “stuck” locks can be easily removed

Non FS Race Conditions

- Linux / BSD kernel ptrace(2) / execve(2) race condition
- ptrace(2)
 - debugging facility
 - used to access other process' registers and memory address space
 - allows to tamper with internal state and execution of a process
 - can only attach to processes of same UID, except when run by root
- execve(2)
 - execute program image
 - setuid functionality (modifying the process EUID)
 - not invoked when process is marked as being traced

execve/ptrace Race

- Problem with execve(2)
 1. first checks whether process is being traced
 2. open image (may block)
 3. allocate memory (may block)
 4. set process EUID according to setuid flags
- Window of vulnerability between step 1 and step 4
 - attacker can attach via ptrace
 - blocking kernel operations allow other user processes to run
- Kernel-side defense against this attack (locking)

Non-FS Race Conditions

- Windows DCOM / RPC vulnerability
- RPCSS service
 - multiple threads process single packet
 - one thread frees memory while other process still works on it
 - can result in memory corruption
 - and thus denial of service

Detection

- Static code analysis
 - specify potentially unsafe patterns and perform pattern matching on source code
 - trivial form:
 - `grep access *.c`
- source code analysis and model checking
 - MOPS (MModel-checking Programs for Security properties)

Detection

- Static code analysis
 - Source code analysis and annotations / rules
 - RacerX (found problems in Linux and commercial software)
 - rccjava (found problems in java.io and java.util)
- Dynamic analysis
 - inferring data races during runtime
 - “Eraser: A Dynamic Data Race Detector for Multithreaded Programs”

Conclusion

- We looked at Race Conditions today
 - Overview
 - TOCTOU
 - Examples
 - Complexity attacks
 - Prevention and Detection